



FALMOUTH
UNIVERSITY

COMP712: Classical Artificial Intelligence

Workshop: Unbeatable AI with Minimax

Dr Daniel Zhang @ Falmouth University
2023-2024 Study Block 1

```
AI's move X(row,col): 0,0
AI's move X(row,col): 2,1
AI's move X(row,col): 1,2
AI's move X(row,col): 2,0
Um, it's a TIE!
```



Table of Contents

- [COMP712: Classical Artificial Intelligence](#)
- [Workshop: Unbeatable AI with Minimax](#)
- [Table of Contents](#)

- [Introduction](#)
- [The Game](#)
- [The Game Tree](#)
- [Strategy Analysis: Minimax](#)
- [Minimax Pseudocode](#)
- [The Repository](#)
- [Your Task](#)
 - [Task 1: Play the demo](#)
- [The Code Structure](#)
 - [Task 2: Implement the minimax algorithm](#)
 - [Task 3: Introduce Improvements](#)
 - [Task 4 \(Optional\): The Nim Game](#)
- [Further Reading](#)

Introduction

[Top](#)

In this session, we'll delve into the core concepts of the `Minimax` algorithm, a fundamental technique in game theory and artificial intelligence. Through practical exercises and discussions, you'll gain a comprehensive understanding of how `Minimax` works and its application in determining optimal strategies in games.

The Game

[Top](#)

We will use a very simple game **Tic-Tac-Toe** to demonstrate how the algorithm works. The rule is: 2 players play in turns on a 3x3 board - one marks **X** and the other marks **O**. The one who gets 3 in a row, a column, or a diagonal first wins.

Since the game board is a 3x3 small grid, the game tree is relatively simple by comparing to other large games like **Gomoku** and **Go**. Therefore, an extensive search of the game tree using the `minimax` algorithm.

The Game Tree

[Top](#)

Source: https://www.researchgate.net/profile/Partha-Sarathi-Chakraborty-2/publication/290786914/figure/fig2/AS:408033621495809@1474294204841/Game-Tree-for-Tic-Tac-Toe_W640.jpg

As shown in the picture above, in this partial representation of the Tic-Tac-Toe game tree, we start from the initial state of the game and showcase a few layers to demonstrate the branching structure.

Observations:

- First Layer:
 - Only 2 nodes are illustrated here, but theoretically, there should be 9 nodes representing all possible moves from the root to reach a state in the first player.
- Subsequent Layers:
 - Each node in the first layer should have 8 child nodes in the second layer, representing the possible moves from the first player.
 - Following this pattern, the third layer should have 7 child nodes for each node in the second layer, and 6 child nodes for each node in the third layer, so on.
- Total Nodes:
 - The total number of nodes in the tree can be calculated as $9 \times 8 \times 7 \times \dots \times 1 + 1 = 362,881$.
 - There might be duplicated nodes due to identical board configurations, but this does not alter the strategic analysis in the next section.

Strategy Analysis: Minimax

[Top](#)

The Minimax algorithm serves as a fundamental approach in decision-making within adversarial games like Tic-Tac-Toe. Its primary objective is to compute the optimal move for a player, taking into account potential moves by the opponent. Operating on the principle of searching through the game tree, Minimax assesses each possible move and assigns a value to nodes representing game states. It utilises a recursive depth-first search (DFS) approach to navigate the tree, evaluating each possible move and its subsequent outcomes. The algorithm assumes rational adversaries, aiming to maximize their advantage while minimizing their opponent's advantage. By simulating hypothetical gameplay,

Minimax identifies the most favourable move for a player, considering all potential sequences of moves by both players until a terminal state or a predefined depth is reached.

Minimax Pseudocode

```
procedure MINIMAX(state, currentPlayer)
  if state == terminal then
    return game over
  else if currentPlayer == maximise then
    initialise bestValue
    for each possible nextState do
      v = MINIMAX(nextState, minimise)
      bestValue = MAX(bestValue, v)
    return bestValue
  else if currentPlayer == minimise then
    initialise bestValue
    for each possible nextState do
      v = MINIMAX(nextState, maximise)
      bestValue = MIN(bestValue, v)
    return bestValue
  end if
end
```

The Repository

[Top](#)

The repository linked below contains the code for the Tic-Tac-Toe game.

Fork the repository (do not clone!) and work on your fork. This will enable you to submit a pull request at the end.

<https://github.falmouth.ac.uk/Daniel-Zhang/COMP712-Mimimax.git>

There are 2 python files in the folder: `demo.pyc` and `minimax.py`. As usual, the `demo.pyc` is a compiled and completed version of the game. And the `minimax.py` is the template file you should be working on in this workshop.

Your Task

[Top](#)

Task 1: Play the demo

[Top](#)

You can play the game by typing the following command:

```
python demo.py
```

Try hard to beat the AI if you can!

The Code Structure

[Top](#)

The code structure is very simply:

- **State** : This class represents the current game state, containing the board and the next player. It includes several member functions facilitating comprehension:
 - **winner** : Decorated as a property of the `State` class, it returns the winner (1 for AI, -1 for human, and 0 for a tie) if the game reaches a terminal state; otherwise, it returns `None`.
 - **move** : Executes a move in the current state and returns the new state after the action, swapping the players.
 - **undo** : Reverses the last move and swaps the players accordingly. A new state is returned after the operation.
 - **is_valid** : Verifies if a given move is valid or not.
 - **get_valid_moves** : Provides a list of all possible moves from the current state.
- Several auxiliary functions manage UI updates and coordinate calculations, which should **remain unchanged**:
 - **draw_line** : Simply draws a line from one point to another.
 - **draw_circle** : Draws a circle centred at the given point with a specified radius.
 - **draw_square** : Renders a filled square at the specified point with a defined side length.
 - **ui_init** : Initializes the board.
 - **ui_update** : Updates the game state.
 - **get_ui_pos** : Converts matrix indices `(row, column)` to UI coordinates `(x, y)`.

- `get_mat_index` : Converts UI coordinates (x, y) to matrix indices (row, column) .
- The last two functions are awaiting completion:
 - `minimax` : This function is the primary algorithm that you'll implement, returning the AI's best move from a given state.
 - `game_play` : This function defines the main interaction of the gameplay.

Note: The mouse click event will be automatically monitored by the turtle module. Therefore, there's no need to concern yourself with human-computer interaction in the `game_play` or `minimax` functions—simply concentrate on the algorithm and game logic.

Task 2: Implement the minimax algorithm

[Top](#)

- You won't be able to play the game until both functions are completed.
- Please refer to [Further Reading](#) or the lecture slides for insights on the `minimax` algorithm.
- Given that the minimax search algorithm possesses complete knowledge of the game tree, a flawless implementation should be unbeatable—the worst outcome it can achieve is a `tie/draw` state.

Task 3: Introduce Improvements

[Top](#)

- Revisit your implementation to explore potential enhancements to halt exhaustive searching like:
 - early-stop
 - alpha-beta pruning
 - or any other methods you come up.
- Play the game and validate your implementation

Task 4 (Optional): The Nim Game

[Top](#)

- Implement an AI agent to play the [Nim Game](#).
- GUI is not necessary; concentrate on the algorithm part.
- Create necessary functions and classes as required.

Note:

You can submit a pull request to the original repository to showcase your work if you like.

Further Reading

- [Online Tic-Tac-Toe Analyser](#)
- [Understanding the Game Tree with Tic Tac Toe](#)
- [Solving Tic-Tac-Toe: Game Tree Basics](#)